(a) Original comparison

FC / FM and same-target context controls

FC | Full clips

Complete input
Complete transcript

FM | Full + fixed-time views

Alternate full clips and 2 s crops
Crop labels include overlapping words

FM-ST | Full-view self-teaching

FM inputs + detached full-view targets
Tests predictive context

CP | Context-preserving views

Restore neighboring input frames
Retain the same crop transcript

(b) Matched supervision

Same source order, seed and 6,000 updates

Full-only

48,000 complete inputs
Complete transcript at every presentation

Fixed-time mixed

24,317 full + 23,683 short views
2 s crops with overlap-based labels

Whole-word mixed

24,317 full + 23,683 short views
2.00–3.28 s word-aligned spans

Whole-word short-only

47,402 short views + 598 full fallbacks
Same word-span rule; no source clip dropped

All states: complete video → unchanged recognizer → complete transcript